

openEuler-24.03 内核KABI白名单删除不必要接口评审

背景：

最近，openEuler-24.03内核回滚mlx5相关模块的CVE/LTS补丁时，发现某些补丁因为mlx5_系列接口在KABI白名单中，会引入kabi变更，比如新增mlx5相关的结构体的成员，我们无法直接适配此类CVE/LTS的修改。mlx5相关驱动社区比较活跃，一般1个月会有5+ bugfix，核心数据结构经常会被修改，最近半年有修改结构体成员的补丁有5+个。

CVE-2024-38556 net/mlx5: Add a timeout to acquire the command queue semaphore

```
854 struct mlx5_cmd_work_ent {
855     unsigned long state;
856     struct mlx5_cmd_msg *in;
857     struct mlx5_cmd_msg *out;
858     void *uout;
859     int uout_size;
860     mlx5_cmd_cbk_t callback;
861     struct delayed_work cb_timeout_work;
862     void *context;
863     int idx;
864     struct completion handling;
865+    struct completion slotted;
866     struct completion done;
867     struct mlx5_cmd *cmd;
868     struct work_struct work;
869     struct mlx5_cmd_layout *lay;
870     int ret;
871     int page_queue;
872     u8 status;
873     u8 token;
874     u64 ts1;
875     u64 ts2;
876     u16 op;
877     bool polling;
878     /* Track the max comp handlers */
879     refcount_t refcnt;
880 };
881
```

Bugfix net/mlx5e: Use the correct lag ports number when creating TISes

```
76 struct mlx5e_resources {
77     struct mlx5e_hw_objs {
78         u32 pdn;
79         struct mlx5_td td;
80         u32 mkey;
81         struct mlx5_sq_bfreg bfreg;
82 #define MLX5_MAX_NUM_TC 8
83         u32 tsn[MLX5_MAX_PORTS]
84         [MLX5_MAX_NUM_TC];
85     } hw_objs;
86     struct net_device *uplink_netdev;
87     struct mutex uplink_netdev_lock;
88     struct mlx5_crypto_dek_priv *dek_priv;
89 };
675 struct mlx5e_resources {
676     struct mlx5e_hw_objs {
677         u32 pdn;
678         struct mlx5_td td;
679         u32 mkey;
680         struct mlx5_sq_bfreg bfreg;
681 #define MLX5_MAX_NUM_TC 8
682         u32 tsn[MLX5_MAX_PORTS]
683         [MLX5_MAX_NUM_TC];
684+        bool tsn_valid;
685     } hw_objs;
686     struct net_device *uplink_netdev;
687     struct mutex uplink_netdev_lock;
688     struct mlx5_crypto_dek_priv *dek_priv;
689 };

```

Bugfix net/mlx5: Fix MTMP register capability offset in MCAM register

```
10309 u8 regs_39_to_32[0x8];
10310
10311+ u8 regs_31_to_11[0x15]; Gal Press
10312 u8 mtmp[0x1];
10313+ u8 regs_9_to_0[0xa];
10314 };
10315
```

mlx5_开头的一系列接口是mlx5驱动内部接口，非内核提供的接口，只与自己的驱动有关，对兼容性影响较小，因此希望删除此类白名单接口。另外，根据此问题，我们又重新排查了下其他白名单接口，并与兼容性团队一起评估过，此列表的接口不必要加入白名单，加入后会导致后续CVE/LTS补丁引入kabi变更无法修复。

兼容性团队刘彦泽 评估意见：

qede、qed、cnic这些驱动在社区不活跃 这些接口属于驱动自身对外开放接口，不属于基础内核开放接口，影响范围较小，只影响本身驱动，对第三方驱动源码兼容安装不受影响 建议可以删除此部分接口。

mlx5_开头接口属于驱动自身对外开放接口，不属于基础内核开放接口，影响范围较小，只影响本身驱动，对第三方驱动源码兼容安装不受影响 建议可以删除此部分接口。

建议方案：

删除openEuler-2403-LTS的kabi 白名单一系列mlx5的接口以及qede、qed与cnic等不必要接口，如此后续CVE/LTS回合与其他代码适配无KABI变更。

待删除接口：

x86白名单接口：

```
qede_rdma_register_driver    drivers/net/ethernet/qlogic/qede/qede
qede_rdma_unregister_driver  drivers/net/ethernet/qlogic/qede/qede
qed_get_fcoe_ops             drivers/net/ethernet/qlogic/qed/qed
qed_get_iscsi_ops            drivers/net/ethernet/qlogic/qed/qed
qed_get_rdma_ops             drivers/net/ethernet/qlogic/qed/qed
qed_put_fcoe_ops             drivers/net/ethernet/qlogic/qed/qed
qed_put_iscsi_ops            drivers/net/ethernet/qlogic/qed/qed
cnic_register_driver         drivers/net/ethernet/broadcom/cnic
cnic_unregister_driver       drivers/net/ethernet/broadcom/cnic
```

arm64白名单接口：

```
mlxfw_firmware_flash        drivers/net/ethernet/mellanox/mlxfw/mlxfw
mlx5_add_flow_rules          drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_alloc_bfreg             drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_blocking_notifier_register
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_blocking_notifier_unregister
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_cmd_check               drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_cmd_cleanup_async_ctx   drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_cmd_create_vport_lag    drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_cmd_destroy_vport_lag   drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_cmd_do                  drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_cmd_exec                drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_cmd_exec_cb             drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_cmd_init_async_ctx      drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_cmd_out_err             drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_comp_eqn_get            drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_comp_vectors_max        drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_access_reg         drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_create_cq          drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_create_mkey        drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_create_psv         drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_create_rq          drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_create_rqt         drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_create_tis         drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_destroy_cq         drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_destroy_mkey       drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_destroy_psv        drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_modify_cq          drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
```

mlx5_core_modify_cq_moderation
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_modify_hca_vport_context
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_modify_rq drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_modify_sq drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_modify_tis drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_query_mkey drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_query_rq drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_query_sq_state drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_query_vendor_id drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_query_vport_counter
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_reserved_gids_count
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_roce_gid_set drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_core_uplink_netdev_event_replay
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_create_auto_grouped_flow_table
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_create_cq drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_create_flow_group drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_create_flow_table drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_create_lag_demux_flow_table
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_db_alloc_node drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_db_free drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_debugfs_get_dev_root drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_debugfs_root drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_debug_qp_add drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_debug_qp_remove drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_del_flow_rules drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_destroy_flow_group drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_destroy_flow_table drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_dm_sw_icm_alloc drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_dm_sw_icm_dealloc drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_eq_create_generic drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_eq_destroy_generic drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_eq_disable drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_eq_enable drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_eq_get_eqe drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_eq_notifier_register drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_eq_notifier_unregister drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_eq_update_ci drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_eswitch_add_send_to_vport_rule
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_eswitch_get_core_dev drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_eswitch_get_encap_mode drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_eswitch_get_proto_dev drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_eswitch_get_total_vports
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_eswitch_get_vport_metadata_for_match
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_eswitch_mode drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_eswitch_register_vport_reps
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_eswitch_unregister_vport_reps
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core

mlx5_eswitch_uplink_get_proto_dev
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_eswitch_vport_match_metadata_enabled
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_fc_create drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_fc_destroy drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_fc_id drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_fc_query drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_fill_page_frag_array drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_flow_table_id drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_frag_buf_alloc_node drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_frag_buf_free drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_free_bfreg drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_get_flow_namespace drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_get_uars_page drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_is_roce_on drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_lag_get_next_peer_mdev drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_lag_get_num_ports drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_lag_get_roce_netdev drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_lag_get_slave_port drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_lag_is_active drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_lag_is_master drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_lag_is_mpesw drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_lag_is_shared_fdb drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_lag_mode_is_hash drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_lag_query_cong_counters
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_modify_header_alloc drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_modify_header_dealloc drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_nic_vport_affiliate_multiport
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_nic_vport_disable_roce drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_nic_vport_enable_roce drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_nic_vport_unaffiliate_multiport
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_nic_vport_update_local_lb
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_notifier_register drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_notifier_unregister drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_packet_reformat_alloc drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_packet_reformat_dealloc
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_put_uars_page drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_qp_debugfs_cleanup drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_qp_debugfs_init drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_query_hca_vport_context
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_query_hca_vport_gid drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_query_hca_vport_node_guid
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_query_hca_vport_pkey drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_query_hca_vport_system_image_guid
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_query_ib_port_oper drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_query_min_inline drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_query_nic_vport_node_guid
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core

```

mlx5_query_nic_vport_qkey_viol_cntr
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_query_nic_vport_system_image_guid
drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_query_port_max_mtu drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_query_port_oper_mtu  drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_query_port_ptys     drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_query_port_vl_hw_cap  drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_rdma_rn_get_params drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_rl_add_rate  drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_rl_add_rate_raw  drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_rl_are_equal  drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_rl_remove_rate drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_rl_remove_rate_raw drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_rsc_dump_cmd_create  drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_rsc_dump_cmd_destroy drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_rsc_dump_next  drivers/net/ethernet/mellanox/mlx5/core/mlx5_core
mlx5_set_port_caps  drivers/net/ethernet/mellanox/mlx5/core/mlx5_core

```

commit 485d65e1357123a697c591a5aeb773994b247ad7

Author: Akiva Goldberger <agoldberger@nvidia.com>

Date: Thu May 9 14:29:50 2024 +0300

net/mlx5: Add a timeout to acquire the command queue semaphore

Prevent forced completion handling on an entry that has not yet been assigned an index, causing an out of bounds access on `idx = -22`. Instead of waiting indefinitely `for` the sem, blocking flow now waits `for` index to be allocated or a sem acquisition timeout before beginning the timer `for` FW completion.

Kernel log example:

```

mlx5_core 0000:06:00.0: wait_func_handle_exec_timeout:1128:(pid 185911):
cmd[-22]: CREATE_UCTX(0xa04) No done completion

```

Fixes: 8e715cd613a1 ("net/mlx5: Set command entry semaphore up once got index free")

Signed-off-by: Akiva Goldberger <agoldberger@nvidia.com>

Reviewed-by: Moshe Shemesh <moshe@nvidia.com>

Signed-off-by: Tariq Toukan <tariqt@nvidia.com>

Link: <https://lore.kernel.org/r/20240509112951.590184-5-tariqt@nvidia.com>

Signed-off-by: Jakub Kicinski <kuba@kernel.org>

diff --git a/drivers/net/ethernet/mellanox/mlx5/core/cmd.c

b/drivers/net/ethernet/mellanox/mlx5/core/cmd.c

index 4957412ff1f6..511e7fee39ac 100644

--- a/drivers/net/ethernet/mellanox/mlx5/core/cmd.c

+++ b/drivers/net/ethernet/mellanox/mlx5/core/cmd.c

```

@@ -969,19 +969,32 @@ static void cmd_work_handler(struct work_struct *work)
     bool poll_cmd = ent->polling;
     struct mlx5_cmd_layout *lay;
     struct mlx5_core_dev *dev;
-    unsigned long cb_timeout;

```

```

-     struct semaphore *sem;
+     unsigned long timeout;
+     unsigned long flags;
+     int alloc_ret;
+     int cmd_mode;

+     complete(&ent->handling);
+
dev = container_of(cmd, struct mlx5_core_dev, cmd);
-     cb_timeout = msecs_to_jiffies(mlx5_tout_ms(dev, CMD));
+     timeout = msecs_to_jiffies(mlx5_tout_ms(dev, CMD));

-     complete(&ent->handling);
-     sem = ent->page_queue ? &cmd->vars.pages_sem : &cmd->vars.sem;
-     down(sem);
+     if (!ent->page_queue) {
+         if (down_timeout(&cmd->vars.sem, timeout)) {
+             mlx5_core_warn(dev, "%s(0x%x) timed out while waiting
for a slot.\n",
+
+                                     mlx5_command_str(ent->op), ent->op);
+             if (ent->callback) {
+                 ent->callback(-EBUSY, ent->context);
+                 mlx5_free_cmd_msg(dev, ent->out);
+                 free_msg(dev, ent->in);
+                 cmd_ent_put(ent);
+             } else {
+                 ent->ret = -EBUSY;
+                 complete(&ent->done);
+             }
+             complete(&ent->slotted);
+             return;
+         }
+         alloc_ret = cmd_alloc_index(cmd, ent);
+         if (alloc_ret < 0) {
+             mlx5_core_err_rl(dev, "failed to allocate command
entry\n");
@@ -994,10 +1007,11 @@ static void cmd_work_handler(struct work_struct *work)
+             ent->ret = -EAGAIN;
+             complete(&ent->done);
+
+         }
-         up(sem);
+         up(&cmd->vars.sem);
+         return;
+     }
+ } else {
+     down(&cmd->vars.pages_sem);
+     ent->idx = cmd->vars.max_reg_cmds;
+     spin_lock_irqsave(&cmd->alloc_lock, flags);
+     clear_bit(ent->idx, &cmd->vars.bitmask);
@@ -1005,6 +1019,8 @@ static void cmd_work_handler(struct work_struct *work)
+     spin_unlock_irqrestore(&cmd->alloc_lock, flags);
+ }

+     complete(&ent->slotted);
+
+     lay = get_inst(cmd, ent->idx);
+     ent->lay = lay;
+     memset(lay, 0, sizeof(*lay));

```

```

@@ -1023,7 +1039,7 @@ static void cmd_work_handler(struct work_struct *work)
    ent->ts1 = ktime_get_ns();
    cmd_mode = cmd->mode;

-    if (ent->callback && schedule_delayed_work(&ent->cb_timeout_work,
cb_timeout))
+    if (ent->callback && schedule_delayed_work(&ent->cb_timeout_work,
timeout))
        cmd_ent_get(ent);
    set_bit(MLX5_CMD_ENT_STATE_PENDING_COMP, &ent->state);

@@ -1143,6 +1159,9 @@ static int wait_func(struct mlx5_core_dev *dev, struct
mlx5_cmd_work_ent *ent)
    ent->ret = -ECANCELED;
    goto out_err;
}

+
+    wait_for_completion(&ent->slotted);
+
    if (cmd->mode == CMD_MODE_POLLING || ent->polling)
        wait_for_completion(&ent->done);
    else if (!wait_for_completion_timeout(&ent->done, timeout))
@@ -1157,6 +1176,9 @@ static int wait_func(struct mlx5_core_dev *dev, struct
mlx5_cmd_work_ent *ent)
    } else if (err == -ECANCELED) {
        mlx5_core_warn(dev, "%s(0x%x) canceled on out of queue
timeout.\n",
                        mlx5_command_str(ent->op), ent->op);
+
    } else if (err == -EBUSY) {
+
        mlx5_core_warn(dev, "%s(0x%x) timeout while waiting for command
semaphore.\n",
                        mlx5_command_str(ent->op), ent->op);
    }
    mlx5_core_dbg(dev, "err %d, delivery status %s(%d)\n",
err, deliv_status_to_str(ent->status), ent->status);
@@ -1208,6 +1230,7 @@ static int mlx5_cmd_invoke(struct mlx5_core_dev *dev,
struct mlx5_cmd_msg *in,
    ent->polling = force_polling;

    init_completion(&ent->handling);
+
    init_completion(&ent->slotted);
    if (!callback)
        init_completion(&ent->done);

@@ -1225,7 +1248,7 @@ static int mlx5_cmd_invoke(struct mlx5_core_dev *dev,
struct mlx5_cmd_msg *in,
    return 0; /* mlx5_cmd_comp_handler() will put(ent) */

    err = wait_func(dev, ent);
-    if (err == -ETIMEDOUT || err == -ECANCELED)
+
+    if (err == -ETIMEDOUT || err == -ECANCELED || err == -EBUSY)
        goto out_free;

    ds = ent->ts2 - ent->ts1;
diff --git a/include/linux/mlx5/driver.h b/include/linux/mlx5/driver.h
index bf9324a31ae9..80452bd98253 100644
--- a/include/linux/mlx5/driver.h
+++ b/include/linux/mlx5/driver.h

```

```
@@ -862,6 +862,7 @@ struct mlx5_cmd_work_ent {  
    void                *context;  
    int                 idx;  
    struct completion   handling;  
+   struct completion   slotted;  
    struct completion   done;  
    struct mlx5_cmd     *cmd;  
    struct work_struct   work;
```